

# Traffic Sign Recognition

## 1. Summary

We proposed a convolution network for German traffic sign classification. It first uses “illumination network” to automatically learn a illumination correction method for preprocessing the input data. By using modified “dense-block”, which are concatenation of convolutions, we are able to “shortcut” activation from bottom to top. Our single network achieve **99.68% accuracy** with just **27.0 millions MAC (multiply accumulation)**.

In order to prevent overfitting of data during training, we use dynamic data argumentation to create new data at the training epoch. This is achieved by randomly applying geometric and illumination transform to the original data.

The code to the project can be found at:

<https://github.com/hengck23-udacity/udacity-driverless-car-nd-p2>

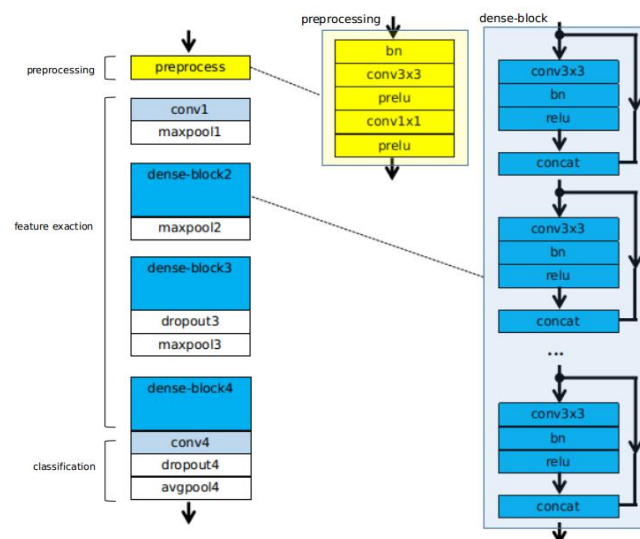


Figure.6: Our network

## 2. Data Set Summary & Exploration

### [Code]

```
In [5]: train_images, train_labels = X_train, y_train

#count
#h = np.histogram(train_labels, bins=np.arange(num_class))

#results image
num_sample=10
results_image = 255.*np.ones(shape=(num_class*height, (num_sample+2+22)*width, channel), dtype=np.float32)
```

### [Description]

The data for the project consist of 32x32x3 rgb cropped image samples of various traffic sign. There are 43 target classes. There are 39,209 samples for training and 12,630 samples, see table.1.

num of samples	train	39209
	test	12630
number of class		43

Table.1: Project data

Figure.1 shows the train data characteristic. The first column is the label image. The second column is the mean image, followed by some example samples of each class. Next is the class label and class name. Finally, we give the sample counts for each class and the class distribution histogram.



Figure.1: Train samples characteristic

We made the following observations:

1. The class distribution is imbalance, with the smallest class having about 200 samples, and the largest with about 2000 samples.
2. Within each class, samples exhibits variations. Brightness and contrast variations seems to be the most prominent. There are also rotation, scale and other geometric transform. Finally there are some blur and minor occlusions in few cases.
3. The mean images for each class represent the class cluster center. They are rather clear and well defined, meaning that the variations are “linear” in the image space. We hence expect the classification problem is “not too difficult”.

Next, Figure.2 shows the test data characteristic. Generally the test data has the same characteristic as the train. Their class distribution is the same. The class examples exhibit the same variations. Most importantly, their mean images are similar.



Figure.2: Test samples characteristic

### 3. Design and Test a Model Architecture

#### 3.1. Setting up of testing, training and validation set

```
[Code] In [10]: #prepare all data here
classnames, X_train, y_train, X_test, y_test = load_data()

train_images, train_labels, valid_images, valid_labels = split_data(X_train, y_train)
test_images, test_labels = X_test, y_test

num_train = len(train_images)
num_valid = len(valid_images)
num_test = len(test_images)
```

#### [Description]

Table.2 shows the setup of the data for testing , validating and training our convolution network. For test data, we use all the 12,630 test samples. For validation, we randomly select 3,000 samples from the original train samples.

<b>test set</b>		12630
<b>validation set</b>		3000
<b>train set</b>	<b>argumented samples per clas</b>	20000
	<b>number of class</b>	43
	<b>total</b>	860000

Table.2: Test, Train and validation data

The remaining  $39,209 - 3,000 = 36,209$  train samples are used to create final augmented set of 86,0000 samples. By using data argumentation, we hope to solve the problems of insufficient train data and class imbalance. The steps of data argumentation are:

1. Flip the 36,209 train samples to create more train samples, see figure.3. This extends the train set to 62,187. We use the flipping method created by <http://navoshta.com/traffic-signs-classification/>
2. Resample the extended set, such that there are 20,000 samples per class. This creates a class-balanced set of  $43 \times 20,000 = 860,000$  samples.
3. Randomly select 80% from the balanced set. These selected samples are perturbed by random geometric transform of rotation, scale, translation and perspective distortion. They are also perturbed by random illumination transform of brightness, contrast and saturation. Figure.4 shows the results of perturbed samples.

#### Flipping

First, we are going to apply a couple of tricks to extend our data by *flipping*. You might have noticed that some traffic signs are invariant to horizontal and/or vertical flipping, which basically means that we can flip an image and it should still be classified as belonging to the same class.



Some signs can be flipped either way — like **Priority Road** or **No Entry** signs.



Figure.4: Example of flipping (adopted from <http://navoshta.com/traffic-signs-classification/>)






Figure.5: Examples of random geometric and illumination transform

### 3.2. Setting up of data preprocessing

**[Code]** - none-

**[Description]**

Popular data preprocessing methods include color transform, whitening and histogram equalization. We do not use such external “hand crafted” data preprocessing. Inspired by paper[1], our network includes data processing layers to act as “illumination transformation network”. Just as “spatial transformer[2]” is being used to automatically correct affine transformation, our “illumination transformer” is used to correct illumination in the input. We will describe the implementation in the next Section3.3 and some of its results in Section3.6.

### 3.3. Setting up of network

**[Code]** In [14]:

```
# my densenet here!
#the inference part (without loss)

def DenseNet_3( input_shape=(1,1,1), output_shape = (1)):

    H, W, C = input_shape
    num_class = output_shape
    input = tf.placeholder(shape=[None, H, W, C], dtype=tf.float32, name='input')
```

#### [Description]

Our network is shown in Table.3 and Figure.6.

\*for max pooling: 2x2/2 means  
\*for dropout: 0.9 means proportion kept

	output			convolution			other*
	H	W	C	num	kernel	mac(M)	
input	32	32	3				
preprocess	32	32	8	8,8	3x3, 1x1	0.3	
conv1	32	32	32	32	5x5	6.6	
maxpool1	16	16	32				2x2/2
dense-block2	16	16	96	concat: 4x16	3x3	8.3	
maxpool2	8	8	96				2x2/2
dense-block3	8	8	192	concat: 4x16	3x3	7.3	
dropout3	8	8	192				0.9
maxpool3	4	4	192				2x2/2
dense-block4	4	4	320	concat: 4x16	3x3	4.4	
conv4	4	4	43		1x1	0.2	
dropout4	4	4	43				0.9
avgpool4	1	1	43				4x4/4
				<b>total</b>		27.0	

Table.3: Our network

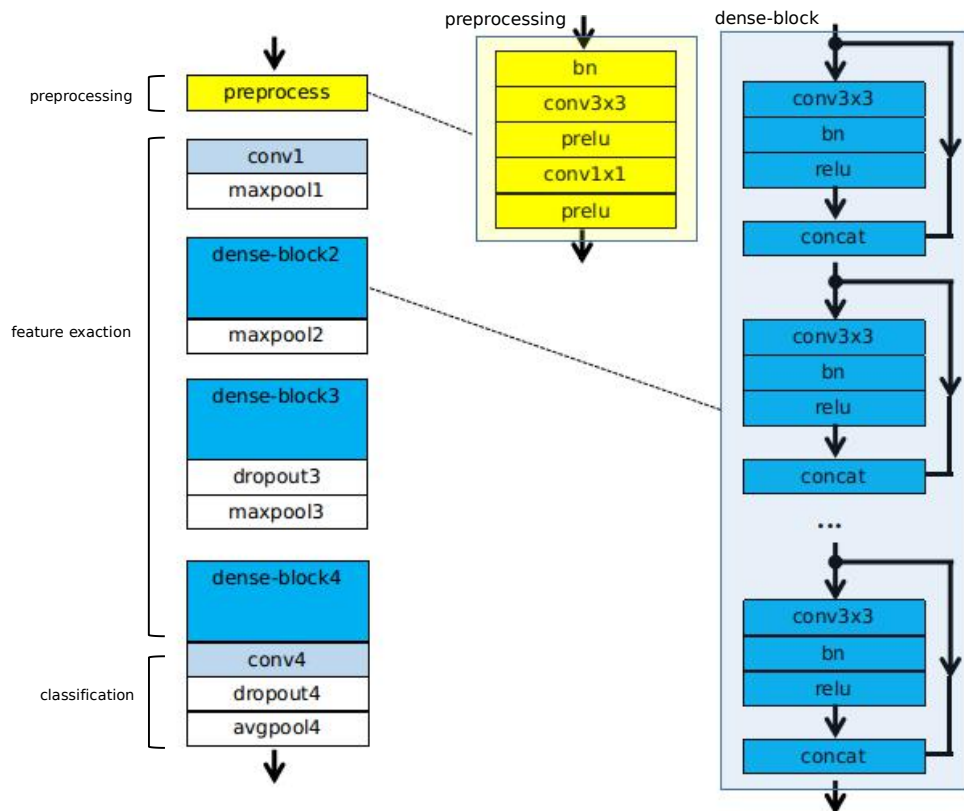


Figure.6: Our network

Our network has total of 27.0 million MACs (multiply-accumulation ops). It has 3 main component:

### 1. Preprocessing

In "preprocess", we first use batch normalization bn layer to standardize the 32x32x3 input to standard normal distribution. We next use 8 conv3x3 later followed by 8 conv1x1, with learnable parametric parametric relu as activation.

### 2. Feature extraction

In "conv1", we use 32 conv5x5, followed by bn and relu activation. Inspired by the work of [3], we next use 3 modified dense-blocks. Each dense-block has 4 concatenation of N conv3x3, followed by bn and relu activation. N =16,24,32 for "dense-block2,3,4" respectively. Max poolings are inserted in between to reduce the activation maps by half successively. The output of the feature extraction is 4x4x320 from "dense-block4".

Note that there are two differences of our dense-block, compared to that of [3]:

- we are using conv-bn-relu-concat, while [3] is using bn-relu-conv-concat
- we do not use dropout within the dense-block. Dropout is only applied after the block if required. In [3], dropout is applied within block, at the conv before concatenation

### 3. Classification

Lastly, in "conv4", the 4x4x320 feature is feed to 43 conv1x1, followed by bn and relu activation. Dropout is used. Global average pooling is applied to the feature to give a logit vector of 1x1x43.

## 3.4. Setting up of solver

```
[Code] In [17]: #solver
epoch_log = 2
max_run = 9
batch_size = 128 #256 #96 384 #128
steps = (0, 3, 6, 8)
rates = (0.1, 0.01, 0.001, 0.0001)

learning_rate = tf.placeholder(tf.float32, shape=[])
solver = tf.train.MomentumOptimizer(learning_rate=learning_rate, momentum=0.9)
solver_step = solver.minimize(loss)
```

### **[Description]**

Our loss function consist of cross entropy loss for classification. L2 regularization loss for the weights is added, with regularization factor=0.0005. We use stochastic gradient descent sgd solver for loss optimization. Batch size is 128. Momentum is set at 0.9.

In order to create an "infinite" number of train samples to prevent overfitting, we dynamically create new argument samples during the training epoch. We think that since the train samples is always changing, it is more difficult for the network to fit the train data. However, we have to be careful that the data change cannot be too big, which else will results in "jumps" in the training loss curve.

Our strategy for dynamic data argumentation is shown in Figure.7. We first divide the epoch into runs. For each run, we select 20% from a fix pool of extended samples, and perturbed the remaining 80% to create new train samples. We then run E epoch of sgd on these train samples. From our experiments, R=9 and E=24 seems to work the best, giving total training epoch of about  $9 \times 24 = 216$ .



Finally, the learning rate is stepped at 0.1, 0.01, 0.001, 0.0001 at the 0,3,6,8 of the runs.

```

▷ Given train samples
▷ Create extended train samples by flipping (steps.1 of Section.3.1)
▷ For run = 1:R
    ▷ Generate new argument samples (steps.2 and 3 of Section.3.1), with
      "new samples = 20% of extended samples + 80% of perturbed samples"
    ▷ For epoch = 1:E
        ▷ Do sgd gradient descent using new samples

```

Figure.6: Our strategy for dynamic data argumentation

### 3.5. Training, validation and testing results

We achieve a good results **99.68% accuracy** and **cross entropy loss 0.012501** on the test set. Table.4 shows the different results. Figure.7 and 8 shows the loss curve and accuracy curve on the train and validation sets. We note that the validation results are better because it does not contain perturbed samples.

train (batch)		valid		test	
loss	acc	loss	acc	loss	acc
0.00142	1.00000	0.00009	1.00000	0.01250	0.99683

Table.4: Training, validation and testing results

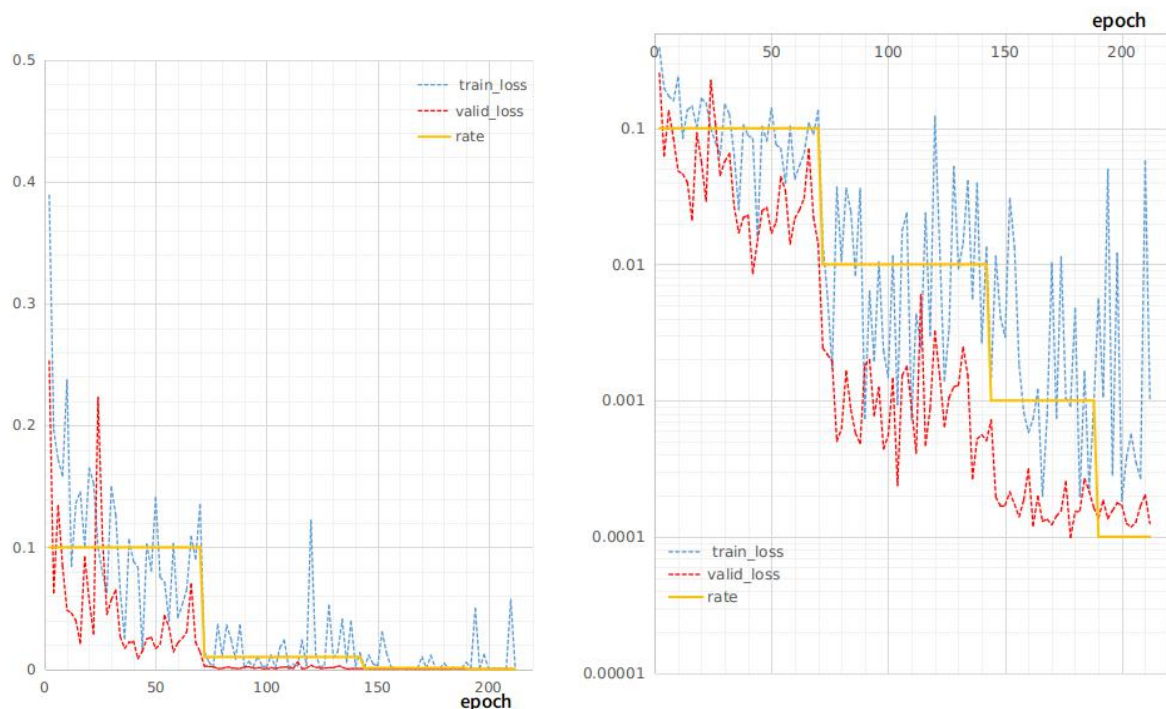


Figure.7: Training, validation loss (right curve is the same left curve in log scale)



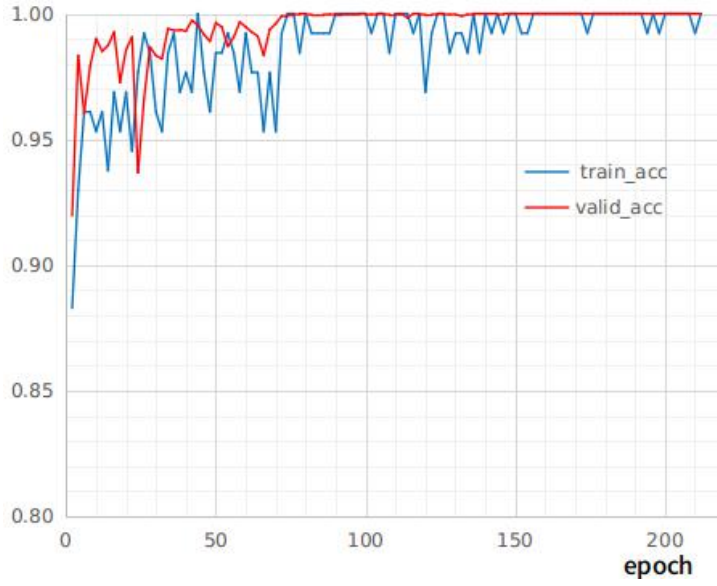


Figure.8: Training, validation accuracy

### 3.6. Discussion of results

#### Results of preprocessing

Figure.9 shows the  $32 \times 32 \times 8$  activation of the preprocessing layer of in Table.3. In (a),(b) and (c), input samples are synthetically generated from our random illumination perturbation, described in section 3.1. (d) are samples from the test set. It can be seen that some of the 8 output channels are sensitive to blue and red colors and invariant to brightness. It is also observed that the few channels has weak activation, meaning that they may be remove for further complexity reduction.

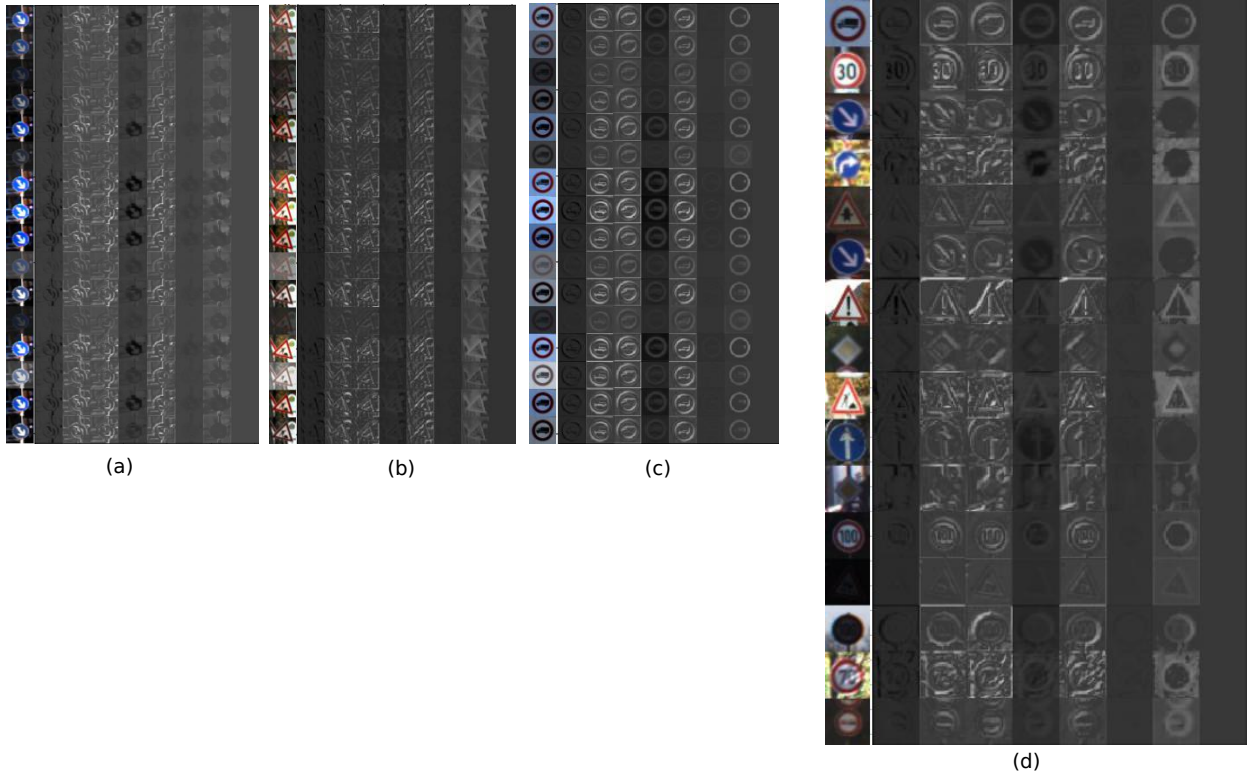


Figure.9: Results of preprocessing on synthetic samples (a),(b),(c) and test samples (d)

### Results of wrong prediction

Figure.10 shows the wrong prediction on test set. There are 40 (0.32%) wrong predictions, possibly due to:

- occlusion
- artifacts on sign (e.g. white paint under the red non-entry sign)
- very small or blurred
- uneven illumination due to sunlight and shadow

Figure.11 shows the expected predicted probability on the test set. Note that this is not the confusion matrix. Given a test sample  $x$  with true label  $y_{\text{hat}}$ , the network predicts  $p_1 \dots p_c \dots p_{43}$  probability for each class. Then:

$$\text{"expected predicted probability}[y_{\text{hat}}, c] = \text{Mean} \{ p_c(x) \}, \text{ over all } x \text{ with true label } y_{\text{hat}}"$$

The figure is to be read row wise. Each row indexes  $y_{\text{hat}}$  and each column indexes  $c$ . (a) is a contrast enhanced version of (b) for better readability. Note that the results of Figure.11 supports the results for Figure.10. For example, "stop sign" as "no-entry sign" is one of the most wrongly predicted results.

### **Network Design Considerations**

We now outline the considerations that lead to the final design of our solution.

#### 1. Establishing baseline performance:

We first do some initial experiments on LeNet and Vgg, see "Appendix.A Additional Experiments on LeNet, Vgg". Note that we do not do a lot fine tuning or extensive hyper-parameter search in these early experiments. We have the following test accuracy:

LeNet Only (MAC = 83.6 M): 89.80%  
LeNet+Flip+Resample: 93.18%  
LeNet+Flip+Resample+Argumentation: 97.61%  
LeNet+Flip+Resample+Argumentation+ dropout: 95.56%  
LeNet+Flip+Resample+Argumentation+ dropout + bn: 97.99%

We conjecture that baseline performance of basic LeNet is about 98%. For Vgg, we can get 98.52% at MAC = 15.5M. These early experiments shows that:

Generating new data is necessary to improve performance  
However, the LeNet network structure has limited accuracy and efficiency. It seems that Vgg can gives better results with less MAC .

Hence we decide to abandon LeNet and design a better structure.

#### 2. Designing a Densenet

From the results of the state-of-art, see "Appendix B : State-of-art performances on traffic sign dataset", we note that multi-scale feature is important. For example, spatial transformer network corrects the scale and Sermanet's multi-scale CNN use skip connections to combine features of low and high scales.

Hence, we choose to use Densenet [3] because it by concatenating conv layers, it can “shortcut” low-scale features to the top. Further, Densenet has shown better performance than other architecture like Vgg, resNet or inception googleNet.

However, we find that Densenet however overfits very easily and results is sensitive to dropout, maybe due to the fact that our problem is small and easy. In fact our initial densenet actually performs worse than Vgg. To reduce overfitting, we make the adjustments:

- Use small number of conv filters

- Remove dropout inside the dense-block (we think due to the small input size of 32x32, the dropout can get magnified by concatenation). Use dropout outside the dense-block instead.

The network design is very much a trial and error process.

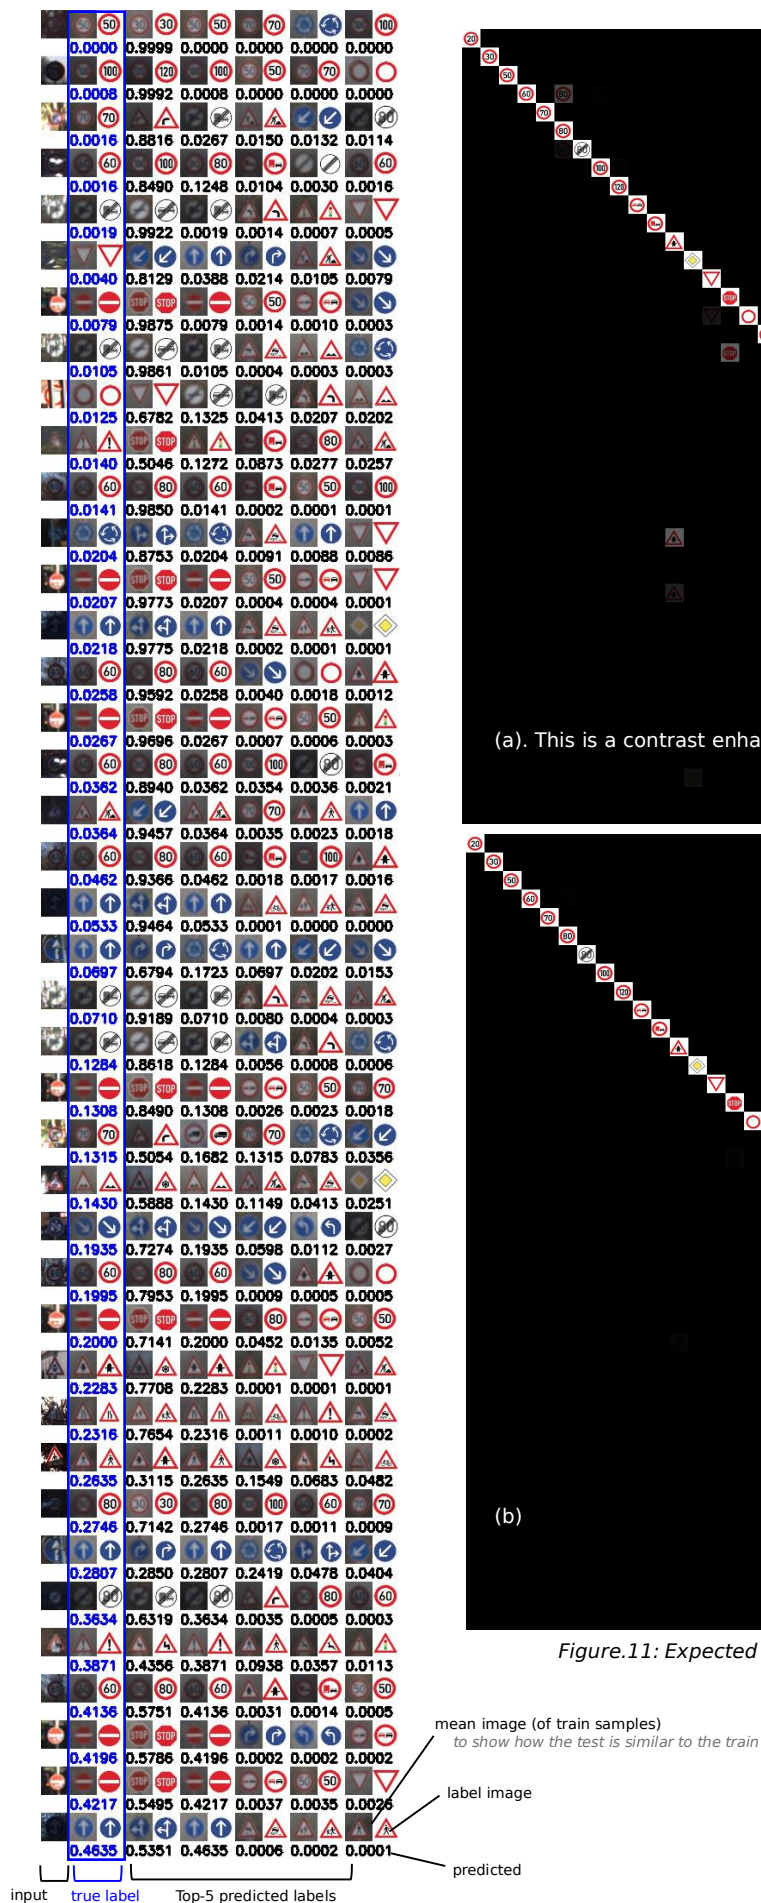


Figure.10: 40(0.32%) wrong predicted test samples




Figure.11: Expected predicted probability (not confusion matrix)



## 4. Testing on New Images

**[Code]** In [20]:

```

### Load the images and plot them here.
### Feel free to use as many code cells as needed.

test_files=['0004.jpg', #normal
            '0000.jpg', #normal
            '0007.jpg', #occluded with snow
            '0006.jpg', #small
            '0005.jpg', #not in class
            ]
test_rois=[(54,180,125,260),(160,430,207,469),(181,32,321,142),(226,65,242,78),(388,408,700,676)]
test_label=[13,38,2,49,-1]

```

### [Description]

Figure.12 shows some test image containing German traffic sign we downloaded from the internet. The roi regions are marked by hand, cropped and resized to 32x32 and input to the trained network from Section.2.











file	true label	crop 32x32x3	original image	difficulty
0004.jpg	13:Yield			Easy.
0000.jpg	38:Keep right			Easy, but sign is off centered in the crop.
0007.jpg	03:Speed limit (60km/h)			Difficult. Sign is occluded by snow. However, human can still recognize it.
0006.jpg	40:Roundabout mandatory			Moderate. However, human may have difficulty in recognizing it in the original image because of the small size
0005.jpg	14: Stop			Difficult? "ARRET" is stop sign in french. This sign is not in the training sample

Figure.12: Test images from the internet

Figure.13 shows the prediction results on the cropped images. Except for image 0007.jpg, we have the rest correct. This is 80% correct. This is lower than the 99.68% of the test set. But the performance may be acceptable, since the wrong prediction is due to heavy occlusion.





file	true label	crop 32x32x3	top-5 predictions (probability label:classname)	comments
0004.jpg	13:Yield		top0: 1.000000 13:Yield top1: 0.000000 26:Traffic signals top2: 0.000000 24:Road narrows on the right top3: 0.000000 00:Speed limit (20km/h) top4: 0.000000 14:Stop	Correct. Confidence is very high
0000.jpg	38:Keep right		top0: 1.000000 38:Keep right top1: 0.000000 34:Turn left ahead top2: 0.000000 05:Speed limit (80km/h) top3: 0.000000 30:Beware of ice/snow top4: 0.000000 36:Go straight or right	Correct. Confidence is very high
0007.jpg	03:Speed limit (60km/h)		top0: 0.352064 28:Children crossing top1: 0.315537 29:Bicycles crossing top2: 0.148701 02:Speed limit (50km/h) top3: 0.037797 19:Dangerous curve to the left top4: 0.025367 40:Roundabout mandatory	Wrong. Confidence is not high at all. The true label is not within top-5. However, the top-3 prediction is speed limit 50 km/h which somehow close to the true sign.
0006.jpg	40:Roundabout mandatory		top0: 0.938421 40:Roundabout mandatory top1: 0.056301 33:Turn right ahead top2: 0.004292 34:Turn left ahead top3: 0.000149 38:Keep right top4: 0.000134 05:Speed limit (80km/h)	Correct. Confidence is very high
0005.jpg	14: Stop		top0: 0.999497 14:Stop top1: 0.000480 29:Bicycles crossing top2: 0.000013 07:Speed limit (100km/h) top3: 0.000003 33:Turn right ahead top4: 0.000002 08:Speed limit (120km/h)	Correct. Confidence is "surprisingly" high. I would expect a correct prediction due to the unique hexagon shape, but I though the confidence is too high.



Figure.13: Prediction of the test images from the internet

## Appendix A : Additional Experiments on LeNet, Vgg

<b>Lenet</b>							
	output			convolution			others
	H	W	C	num	kernel	mac(M)	
input	32	32	3				
conv	32	32	108	108	5x5	8.3	
maxpool	16	16	108				2x2/2
conv	16	16	108	108	5x5	74.6	
maxpool	8	8	108				2x2/2
dense			100	100		0.7	
(droupout)							0.5
dense			100	100			
(droupout)							0.5
dense			43	43			
<b>total</b>						83.6	

<b>Vgg</b>							
	output			convolution			others
	H	W	C	num	kernel	mac(M)	
input	32	32	3				
conv	32	32	64	64	5x5	4.9	
maxpool	16	16	64				2x2/2
conv	16	16	32	32	1x1	0.5	
conv	16	16	32	32	3x3	2.4	
conv	16	16	64	64	1x1	0.5	
maxpool	8	8	64				2x2/2
conv	8	8	64	64	1x1	0.3	
conv	8	8	64	64	3x3	2.4	
conv	8	8	128	128	1x1	0.5	
conv	8	8	64	64	1x1	0.5	
conv	8	8	64	64	3x3	2.4	
conv	8	8	128	128	1x1	0.5	
maxpool	4	4	128				2x2/2
dense			256	256		0.5	
dense			256	256		0.1	
dense			43	43		0	
<b>total</b>						15.5	

Figure.A.1: LetNet and Vgg network used in additional experiments

comparing preprocessing				argumentation			train (batch)		valid		test	
		MAC (M)	preprocess	flip + resample	perturbance geometric illumination		loss	acc	loss	acc	loss	acc
Lenet_2	lenet + dropout + bn	83.6	rgb=(rgb-128)/128	•	•		0.1274	0.9531	0.0024	0.9997	0.0852	0.9779
Lenet_3	lenet + dropout + bn + whiten			•	•		0.1448	0.9609	0.0129	0.9997	0.0959	0.9698

comparing augmentation				argumentation									
				flip + resample	perturbance		train (batch)		valid		test		
		MAC (M)	preprocess		geometric	illumination	loss	acc	loss	acc	loss	acc	
Lenet_0	lenet	83.6	rgb=(rgb-128)/128				0.0227	1.0000	0.0957	0.9770	0.5318	0.8980	
				●			0.0000	1.0000	0.0669	0.9903	0.6316	0.9318	
				●	●		0.0218	0.9922	0.0052	0.9990	0.2713	0.9561	
				●		●	0.0372	0.9922	0.0729	0.9913	0.5969	0.9512	
				●	●	●	0.1615	0.9531	0.0043	0.9980	0.1338	0.9761	
Lenet_1	lenet + dropout			●	●		0.2380	0.9375	0.0273	0.9933	0.1723	0.9682	
				●		●	0.5629	0.8594	0.0512	0.9867	0.1621	0.9556	
				●	●		0.1274	0.9531	0.0024	0.9997	0.0852	0.9779	
Lenet_2	lenet + dropout + bn			●		●	0.0504	0.9844	0.0060	0.9983	0.0832	0.9817	
				●	●	●	0.2666	0.9297	0.0037	0.9987	0.0730	0.9799	
				●	●		0.0002	1.0000	0.0022	0.9993	0.1143	0.9721	
Vgg_0	vgg + bn	15.5		●		●	0.0332	0.9922	0.0112	0.9970	0.1818	0.9599	
				●	●	●	0.1109	0.9688	0.0027	0.9997	0.0624	0.9852	

Figure.A.2: Training and validation results in additional experiments




## Appendix B : State-of-art performances on traffic sign dataset

	test accuracy
Industrial performance/Cadence hierarchical CNN [3]	99.82
<b>illumination Transformer + DenseNet (Ours)</b>	<b>99.68</b>
2 Spatial Transformer Networks[8]	99.67
Hinge Loss CNN[5]	99.65
Spatial Transformer Networks[4]	99.61
IDSIA Committee of CNNs[2]	99.46
Alex modified multi-scale CNN [6]	99.33
Vivek modified Vgg CNN [7]	99.10
Sermanet multi-scale CNN [1]	99.17
Human performance[1]	98.84

- [1] "Traffic Sign Recognition with Multi-Scale Convolutional Networks" - Pierre Sermanet and Yann LeCun, IJCNN 2011  
[2] "Multi-column deep neural network for traffic sign classification" - D. Cireşan, U. Meier, J. Masci, and J. Schmidhuber, NN 2012.  
[3] "Is Bigger CNN Better?" - Samer Hicazi - eNNS2016  
[https://ip.cadence.com/uploads/presentations/1345PM\\_ENNS\\_v10\\_Samer\\_Hicazi.pdf](https://ip.cadence.com/uploads/presentations/1345PM_ENNS_v10_Samer_Hicazi.pdf)  
[4] [http://torch.ch/blog/2015/09/07/spatial\\_transformers.html](http://torch.ch/blog/2015/09/07/spatial_transformers.html)  
[5] "Traffic Sign Recognition With Hinge Loss Trained Convolutional Neural Networks" - Junqi Jin, Kun Fu, Changshui Zhanq, IEEE ITS 2014  
[6] <http://navoshta.com/traffic-signs-classification/>  
[7] <https://medium.com/@vivek.yadav/improved-performance-of-deep-learning-neural-network-models-on-traffic-sign-classification-using-6355346da2dc#.2c5sqna28>  
[8] [https://github.com/Moodstocks/qtsrb\\_torch](https://github.com/Moodstocks/qtsrb_torch)

Figure.B.1: Comparing performances (accuracy)



[https://community.cadence.com/cadence\\_blogs\\_8/b/breakfast-bytes/archive/2016/07/13/how-to-optimize-your-cnn](https://community.cadence.com/cadence_blogs_8/b/breakfast-bytes/archive/2016/07/13/how-to-optimize-your-cnn)

Figure.B.2: Comparing performances (MACs and accuracy)



## References

- [1] "Systematic evaluation of CNN advances on the ImageNet"-Dmytro Mishkin, Nikolay Sergievskiy, Jiri Matas, Arxiv 2016
- [2] "Spatial Transformer Networks” - Max Jaderberg, Karen Simonyan, Andrew Zisserman, Koray Kavukcuoglu, Arxiv 2015
- [3] "Densely Connected Convolutional Networks" - Gao Huang, Zhuang Liu, Kilian Weinberger, Laurens van der Maaten, Arxiv 2016